

## 1 Function approximation

- (a) (Exercise 9.1 from Sutton): Show that tabular methods such as presented in Part I of Sutton's book are a special case of linear function approximation. What would the feature vectors be?

## 2 Feature designing




Figure 1: Super Mario Bros., ©NINTENDO 1985

- (a) Imagine you want to implement a RL agent to play Super Mario (the original one). For the policy  $\mu(x)$  to work, we need to construct features. Please give a example feature vector and what each component encodes.
- Hint:** A feature vector is given by  $x = \langle x_1, x_2, \dots \rangle$ , where  $x_1$ , for example, encodes the x-coordinate of Mario,  $x_2$  encodes the y-coordinate of Mario, etc. Please give a feature vector of at least 5 dimensions.

### 3 Value function fitting

In this exercise, you will fit an approximated value function for a given policy. `valueapprox.zip` contains the following files:

**custompendulumenv.py:** Custom gym pendulum environment. This is a classic environment, Mujoco is not needed. (no changes required)

**Note:** Instead of a rewards, the environment returns a cost that is zero at the optimum and positive away from the optimum. This means that lower values are better.

**agent.py** This file implements the agent class. For you, the agent is a black box<sup>1</sup> that exposes a method `get_action` which you can use to query an action  $u$  for a given state  $s$ . (no changes required)

**memory.py** Implements a replay buffer that provides a method `addTuple` to add new transitions to the buffer and `sample` that, in our case, returns a tuple of the form  $(s_t, u_t, c_t, s_{t+1})$  where  $s_t$  is the current and  $s_{t+1}$  the next state,  $u_t$  is the action and  $c_t$  the cost. (no changes required)

**pendulum\_value\_plot.ipynb** In this notebook, you will implement the value fitting with a function approximator (NN). The notebook provides already routines for data collection and value function visualization. (Changes required)

- (a) Use a function approximator, like a NN, to fit the value function. Use an appropriate target (td update) for the value fitting.
- (b) Give an intuitive interpretation of the plotted value function. In case you could not solve part one of the exercise, Figure 2 shows a plot of the expected output.

---

<sup>1</sup>In fact, the agent outputs a linear mapping  $u = -Ks$  from inputs  $s$  to outputs  $u$ . The gain factor  $K$  is derived from a linear quadratic regulator (LQR)




Figure 2: Approximated value function for a linear policy and the pendulum environment. Lower values are better.